


Inverted tokenization: one token per variate, not per timestep



Attention runs over N variates, never over L timesteps, so there is no causal mask: causality is enforced upstream, in the features and the windowing (root section 6.1). Loss is MSE on the TARGET CHANNEL ONLY at every rung, so auxiliary supervision does not vary with K , which is the study's own independent variable (D39).